

ENGAGEMENT PROPOSAL · BUILD TO LAUNCH & PHASE 3

The agentic layer for the TXN platform.

From a six-week agentic pilot to version one of every component across the Console, the Developer Portal, and TXN's own operations. Built to ship the differentiation before the window closes.

Prepared for
Ian Johnson, CEO, TXN

Prepared by
Novosapien

Date
19 July 2026

Version
1.0 · Draft

01

Executive summary

TXN is building the first agentic-native issuer processor. Sprint Zero is delivered and the requirements are aligned. This proposal sets out how Novosapien builds the agentic layer that carries that bet: a working agentic pilot in six weeks, then the launch-critical components, including the full Agentic experience in production, wired in for the October launch. Phase 3 takes the remaining deliverables to version one across a committed three months, November to January, with the team retained monthly at the same rate.

6 weeks

to a working agentic pilot, on simulation, no vendor dependency

4 months

to the October launch, with the full Agentic experience live in production

8

components, taken to version one across the engagement

I don't think we can do concept one and then at some point in the future do concept two, because somebody will beat us to it.

Ian Johnson, CEO, TXN

The plan answers that concern directly. We do not sequence the agentic experience behind a year of platform plumbing. The pilot proves Concept 2 and Concept 3 in six weeks on simulated tools, so the differentiation is demonstrable while Direct Transact and Stackworkz finish the surfaces it will later run on. As those dependencies land, we wire the same agentic layer into the real APIs, data lake, and Console.

What this proposal covers

The remaining build to the October launch, a three-month Phase 3 (November to January) of remaining deliverables to version one with scope agreed with TXN, the flight plan and vendor dependencies, the delivery team, and the full commercial model.

What stays with TXN

The Core API and data lake (Direct Transact), the Console and Developer Portal (Stackworkz, design by Super Ultra), all compliance frameworks, and pricing strategy. Novosapien delivers the agentic layer only.

02

Where we are today

This is not a standing start. Sprint Zero is complete and the engagement is live. This proposal continues work already underway, now delivered under the Novosapien name following the management buyout.

Delivered in Sprint Zero COMPLETE

Requirements alignment

A validated, prioritised backlog with acceptance criteria across the full component set, agreed with TXN's stakeholders.

Architecture and approach

The solution architecture, the model-tier routing and cost-governance approach, and a verified LLM operating-cost model across the component set.

Partner and dependency mapping

A dependency map across Direct Transact, Stackworkz, and Super Ultra, with the integration points the agentic layer needs from each.

Component deep-dives

Scope, sub-components, and sequencing captured for the Agent Access Layer, Alerts & Inbox, Co-pilot, Full Agentic Experience, Developer Support, and Internal Ops.


Commercial position. Sprint Zero (£17,500) is delivered, with the invoice being issued this week under Novosapien. This proposal allocates the remaining build budget and proposes the retainer extension that takes the platform from pilot to a complete version-one agentic layer.

The eight components in scope

Component	Role	Status	V1 Build Size
Agent Access Layer (+ A2A endpoint)	Foundational tool surface every agent rides on	DEFINED	L
Co-pilot	Reactive in-console assistant (Concept 1)	DEFINED	M
Agent Inbox & Alerts	Proactive analysis and approvable plans (C1 to C2)	DEFINED	M
Full Agentic Experience	Agent as the interface (C2 to C3)	DEFINED	XL
Developer Support	Docs MCP server, portal co-pilot, support triage	DEFINED	S
Internal Ops Agents	Run TXN agentially, onboarding first	DEFINING	XL
Prepared by Novosapien for TXN			
Fraud & Risk Assist	Real-time enrichment and rules, advise not decide	COLLECTING	M
Commercial in confidence · 02			

02 Solution architecture

How the parts connect. The Novosapien agentic layer (Co-pilot, Agent Inbox, Full Agentic Experience, Developer Support, Internal Ops, Fraud & Risk, and Reconciliation) renders inside TXN's surfaces and rides on a single Agent Access Layer that exposes the Core API as permission-scoped tools. The A2A endpoint is the inbound door for a client's own agent.



■ Novosapien delivers
 □ Direct Transact / TXN
 ⋈ Stackworkz surface (agent renders inside)
 □ External client agent

The bet, and why timing decides it

TXN's differentiation is not feature parity with Marqeta or Galileo. It is being agentic-native where they are retrofitted. That is a structural advantage an incumbent cannot copy quickly: they cannot redesign a console built for expert users into an agent-first product without disrupting the customers they already have.

The incumbent model

- API-first but operationally heavy: clients still need card experts to configure and run programs.
- 50 to 60 configuration properties exposed raw, with no business-language wrapper. People turn them on and off and break things.
- AI bolted onto a console designed for experts. A side panel, not the interface.

The agentic-native model

- The user does not become a card expert. The agent is the card expert.
- Onboarding collapses from a multi-week exercise into a guided conversation.
- The agent is the primary interface from day one, with a human in the loop until trust is earned.

I want to know as little as possible about the details of card programs. I don't want to employ a card expert. I want you to be my card expert.

Ian Johnson, CEO, TXN

The agentic stack that makes this buildable, MCP, A2A, dependable tool-use, multi-step reasoning, reached production-readiness only recently. A small lead in agentic maturity converts into a structural product advantage during this window. The cost of waiting is concrete: ship Concept 1 alone and the product can feel a year behind on launch day.

The plan is built around that risk. The pilot demonstrates the agentic experience in six weeks, independent of the other vendors' timelines, so TXN can show the market its difference now rather than after the platform plumbing is finished.

How we deliver, and how we de-risk it

Every component graduates along the same trust ladder, mapped to TXN's three concepts. A human stays in the loop until the agent has earned autonomy, and the platform's approval queues are respected at every step.

CONCEPT 1

Co-pilot

The user drives. The agent recommends and explains impact before anything changes. Human-initiated, human-validated.

CONCEPT 2

Agent-advisor

The agent initiates: "I have noticed this, here is the plan, may I execute?" The named approver signs off. Agent-initiated, human-validated.

CONCEPT 3

Full agentic / A2A

The client's own agent talks to TXN's agent through MCP and A2A, with the same permissions and approvals as the human it represents.

Map, Train, Deploy, Scale

Map

Lock the work each component owns, its tools, permissions, and the approval gates it must respect. Largely delivered in Sprint Zero.

Train

Ground each agent in TXN's documentation and the live API contract, then harden against the failure modes that matter in payments.

Deploy

Ship behind feature flags and a kill switch, starting human-validated, with full audit on every agent action.

Scale

Earn autonomy component by component as the data flywheel fills and confidence is evidenced in production.

Governance is built in, not bolted on. Permission scoping mirrors the Console, high-blast-radius changes route through the existing two-person approval queue, every agent action is logged to an append-only audit trail (immutable by design; storage being finalised with DT), and model-tier routing plus caching and budget thresholds keep consumption bounded by design.

The agentic layer

Eight components, organised by how the agent shows up, not by technology. Three client-facing lanes along the trust ladder, a foundational tool surface beneath them, the developer surface, and the agents that run TXN itself.

CLIENT-FACING · RENDERED INSIDE THE CONSOLE

Co-pilot

M

Context-aware chat, guided configuration, impact preview before any change. Concept 1.

Agent Inbox & Alerts

M

Change-impact, threshold, and scheduled analysis surfaced as actionable, approvable plans. C1 to C2.

Full Agentic Experience

XL

The agent as interface: renders live UI, composes journeys, acts on approval. C2 to C3.

FOUNDATIONAL · EVERY AGENT RIDES ON THIS

Agent Access Layer (+ A2A endpoint)

L

Wraps the Core API as permission-scoped, agent-callable tools, exposed via MCP. Three-layer permission validation, approval-as-permission on high-risk changes, and combined Console, API, and chat audit. The A2A endpoint is its external edge: the inbound door for clients' own agents.

DEVELOPER PORTAL

Developer Support

S

A hosted docs MCP server (the strategic surface: developers plug it into their own agents), a light portal co-pilot, sandbox assist, and triage that packages well-formed tickets into Internal Ops.

PAYMENT STREAM · LATER PHASE

Fraud & Risk Assist

M

Real-time enrichment of the approve/decline pass-through plus a rules engine and recommendations. Advise, don't decide. Data-dependent, sequenced once transaction volume accumulates.

SETTLEMENT · DUAL-USE

Reconciliation

TBC

Auth and clearing settlement reconciliation. TXN uses it internally to verify the platform; clients use it heavily too. Promoted from Internal Ops on 18 June; scope confirmed in a capture session with DT.

INTERNAL · RUNS TXN AGENTICALLY

Internal Ops Agents

XL

Customer onboarding (the first build target), a self-healing knowledge engine, support triage, the release pipeline, and a simulation harness for the cold-start period. CRM is the source of truth. Keeps headcount flat against client growth.

Internal operations are part of the product, not a side-project. The same agentic philosophy that powers the client surfaces powers the business that runs them, so the experience clients see is not throttled by humans doing manual work behind it.

Phase 1: the six-week agentic pilot

The pilot exists to prove the experience that TXN is betting on, before the other vendors' surfaces are ready. It runs on simulated tools and synthetic data, so almost nothing in it waits on Direct Transact or Stackworkz.

What we build

- A working **Full Agentic Experience** showcase: a Claude-like interface that takes a goal, plans, renders live UI in the conversation, and acts on approval.
- The **Agent Access Layer** foundations: the MCP tool surface and permission model, running against mocked Core API endpoints built from the YAML and sample payloads.
- **Internal Ops v1**: the customer-onboarding agent, meeting capture, and the knowledge engine, the low-dependency, CRM-based parts that return value from day one.

Why pilot-first

- It demonstrates Concept 2 and Concept 3 now, answering the "someone will beat us to it" risk head-on.
- It de-risks the hardest component (the XL Full Agentic Experience) on simulation before a penny of wire-in spend.
- It surfaces gaps in the API contract and the data the agent needs, feeding the wire-in phase.
- Internal Ops v1 needs no DT or Console dependency, so it banks value while the platform stabilises.

Outcome of Phase 1. A demonstrable agentic experience TXN can show to its market and its board, plus a hardened foundation (tool surface, permission model, simulation harness) that the rest of the build plugs into.

Wks 1-6

20 July to end-August 2026

0

vendor dependencies on the critical path

3

components advanced: Full Agentic, Access Layer, Internal Ops v1

Phase 2: what we deliver to launch

Phase 2 wires the agentic layer into production through the four months to the October launch, set out component by component. Each is built synthetic-first: against mock APIs and data, with the demo-data flags dropping as the real surfaces land, so nothing stalls waiting on the platform. This page covers the launch-critical layer; the next covers the Developer Portal and Internal Ops.

Agent Access Layer L the foundation every agent rides on

- A single MCP tool surface that wraps the Core API as business-language tools (issue or suspend a card, configure a product, query transactions), gated per user, permission, and turn.
- A permission model that mirrors the Console's granular roles, plus an AI-specific permission config, with approvals routed into the existing Console approval queue under the two-person rule.
- A combined audit trail across Console, Core API, and AI chat, so every agent action is attributable and reconstructable.
- The A2A endpoint: an inbound door for a client's own external agent, reusing the same tools, scoping, approval, and audit.

Full Agentic Experience XL the showcase, agent as the interface

- A Claude-like surface: state a goal in natural language, the agent plans, renders live interface inside the conversation, and acts only on approval.
- Generative UI: the agent calls a tool and the front end renders the matching component from the Console library, live and clickable, deep-linking back into the Console.
- Multi-agent planning that composes across journeys (for example, "change a transaction limit and show me the impact"), runs steps concurrently, and executes in order on approval.
- Saved, revisitable views and scheduled tasks, proven on roughly 100-run guard-rail simulations, then wired onto the live Core API so the full experience runs in production before mid-October.

Co-pilot v1 M reactive in-Console assistant

- Natural-language Q&A and navigation, page-scoped with page context and guard rails.
- Guided configuration that collapses 100+ product-config fields into a confirm-not-configure conversation, surfacing what most clients in a given pattern do.
- Action on confirmation: plan-then-act for multi-step flows via the Access Layer, with explicit button confirmation and bundled approvals.

Phase 2: what we deliver to launch, continued

The Developer Portal lands at the end of July. Internal Ops returns value from day one because it needs no transaction data.

Alerts v1 M proactive change-impact and scheduled alerts

- Change-impact alerts the moment a setting changes, whatever the channel (Console, API, or AI).
- User-defined threshold monitors on a standard observability platform, with the AI invoked only when a threshold trips, to control cost.
- Scheduled analysis and reporting: cadence summaries with the drivers behind the numbers, on-demand reports, and reusable templates.

Developer Support S for the end-July Developer Portal

- A hosted docs MCP server developers plug straight into Claude or Claude Code, API-key-gated and rate-limited.
- A portal co-pilot that answers from the docs corpus, grounded and deterministic, surfacing known answers and resolved tickets.
- Support triage that reads the developer's recent API logs, diagnoses the likely cause, and packages a well-formed ticket the agent swarm pre-investigates before a human sees it.
- Four access levels (visitor, signed-up, prospect, client) that double as lead capture.

Internal Ops v1 XL value from day one, no transaction data needed

- Customer onboarding: automated due diligence with fully sourced reports (Companies House and data providers), an onboarding-interview agent that captures intent and surfaces hard realities (the roughly 12-week Visa lead time), Statement of Work assembly with drift detection, and CIQ assembly for the processor-side fields.
- Meeting capture: a notetaker records the call, looks attendees up in the CRM, routes the transcript to the right analysis team, and writes structured notes and actions back.
- Knowledge engine: one central knowledge base (Umbraco) that improves itself three ways, reactive capture from support, proactive mining of recurring questions, and self-healing docs (a production error opens a fix PR), every change human-approved.

Built synthetic-first. Every component is built against mock APIs and simulated surfaces, validated in simulation, then switched to the real Core API, data lake, and Console as each lands. Progress does not wait on any single vendor dependency.

9 Jul

Prepared by Novosapient for TXN

Console handover (Stackworkz design pack)

end-Jul

Developer Portal live: docs MCP, portal co-pilot, triage


early-Oct

Confidential in confidence · 06

Market launch: Console + Co-pilot, Full Agentic, Alerts v1 live

Flight plan

Three phases beyond Sprint Zero. The build phase commences around 20 July and carries the launch-critical components, including the full Agentic experience wired into production, toward the October market launch. Phase 3 then takes the remaining deliverables to version one over a committed three months, November to the end of January, with the tail landing in early February, on a scope agreed with TXN. We build synthetic-first: each component is built against mock data and APIs in parallel, then the demo-data flags drop as the real APIs and data land, so progress does not stall waiting on the platform.



PHASE	WINDOW	DELIVERS	GATED BY
0 · Sprint Zero	Complete	Requirements, architecture, dependency map	Nothing
1 · Agentic pilot	Wks 1-6 · from 20 Jul	Full Agentic showcase, Access Layer foundations, and Internal Ops v1 (onboarding, capture, knowledge)	LOW CRM-based, runs on simulation
2 · Wire-in	Aug to mid-Nov	Developer Support for the end-July portal; Access Layer on the real API; the full Agentic experience wired into production; Co-pilot v1 and Alerts v1 for the October Console launch	DT API, STACKWORKZ, INFRA SEPARATION
3 · Remaining deliverables	Nov to early Feb · 3 months	The remaining deliverables taken to version one, scope agreed with TXN	DATA LAKE, VOLUME

Two fixed dates The Developer Portal launches end-July and the market launches early October (the partners' stated targets). With the build commencing around 20 July, Developer Support fast-follows the portal, and the in-Console components wire in from the 9 July Console handover toward the October launch.

Dependencies

Novosapien plugs into a multi-vendor build. The pilot is designed to need none of it. Everything after the pilot has a clear owner and a clear readiness gate.

PROVIDED BY	WHAT THE AGENTIC LAYER NEEDS	GATES WHICH WORK
Direct Transact Core API + Data Lake	Stable API endpoints and the YAML contract; data-lake schema and access for analytics, alerts, and audit storage	Access Layer wire-in, Co-pilot data reads, Inbox anomaly scans, Fraud & Risk, audit trail
Stackworkz Console + BFF + Portal	Console instrumentation (page state, component handles, action handlers); BFF permission model (the source of truth the agent mirrors); Umbraco docs/guides API (reference sourced from the DT YAML)	Co-pilot and Inbox in-console, Full Agentic generative UI, Developer Support docs
Super Ultra Console design	The component library and interaction patterns the agent renders into	Full Agentic Experience generative UI
TXN Product + data	CRM (Freshsales) as source of truth; AI permission config; brand guidelines; transaction volume over time	Internal Ops onboarding, permission scoping, Fraud & Risk maturity
Providers Anthropic / Google	A thin MCP-as-message passthrough is the confirmed A2A design; direct provider A2A is a later enhancement	External client agents (A2A)

Open items that affect scope or sequence

API readiness

Around 50 endpoints in the current YAML; functional testing is underway (Postman and automation). Endpoint names and fields are still settling; the architecture is stable. This sets mock-API fidelity for the pilot and wire-in.

Infrastructure separation

TXN-to-DT separation: DT's build must be stood up in TXN's own Azure (Europe) before we can wire to the real API. The current critical-path dependency. MCP ownership is now settled (Stackworkz docs MCP, DT card-API MCP).

Data retention location

Data residency is settled (EU isolation, with prompt and log minimisation and redaction). The seven-year PCI retention location is the remaining open item with DT.

Open with DT

Reconciliation is now its own component (dual-use). Its scope, the sandbox API-key model (JWT versus long-lived), product-webhook timing (post-Visa-certification), and alerting ownership are being finalised with DT.

Delivery team and ways of working

A dedicated, full-time team of three carries the build and then the three-month Phase 3 at the same monthly fee. The same continuity from pilot to version one is what keeps the agentic layer coherent across surfaces.

AI Product Owner

Full-time. Owns scope, sequencing, acceptance, and the line to TXN's stakeholders and the other vendors. Holds the agentic vision across components.

AI Engineer

Full-time. Builds the agents, the MCP tool surface, the orchestration, evaluation, and guardrails.

Forward Deploy Engineer

Full-time. Wires the layer into the real APIs, Console, data lake, and Developer Portal as each dependency lands.

Ways of working

Cadence

Fortnightly demonstrations of working software to TXN stakeholders, with a standing weekly management update.

Quality

An evaluation harness with regression suites validates every prompt and agent change before it ships. Agent observability traces calls, latency, tokens, and quality.

One team, end to end. The same three people run the four-month build to launch, then continue through the three-month Phase 3 to January. No re-mobilisation, no knowledge loss, no ramp tax between phases.

Standard day rates

ROLE	STANDARD DAY RATE
AI Product Owner	£950
AI Engineer	£850
Forward Deploy Engineer	£750

The retainer is the value. The retainer secures all three roles full-time for £36,125 per month. At standard day rates the same team is £2,550 per day, about £53,550 per month across 21 working days, so the retainer runs roughly a third below buying the equivalent days. Change requests outside the retainer are charged at these standard rates.

Investment

The remaining build budget to launch, the Phase 3 deliverables that remain by component, the three-month team retainer to January, and a fixed support and maintenance fee. All figures exclusive of VAT.

A · Build phase to launch

Component Focus	Allocation
Full Agentic Experience (production, wired)	£48,000
Agent Access Layer (foundation, mock to real API)	£36,500
Co-pilot v1	£18,000
Developer Support (for the end-July portal)	£15,000
Alerts v1	£12,000
Internal Ops v1 (onboarding, capture, knowledge)	£15,000
Build phase total · 4 months	£144,500

Remaining budget of £162,000 less the £17,500 Sprint Zero, delivered across four months by the full-time team. Allocations are indicative of where capacity lands, not fixed price-per-component.

B · Phase 3, the remaining deliverables

After the October launch, Phase 3 runs as a committed three-month engagement across November, December, and January, taking the remaining deliverables to version one. They are set out in detail on the next two pages, component by component. TXN chooses which deliverables are in scope and in what order; the team works them in parallel, and a substantial share of the menu is deliverable within the three months. If kickoff moves, the window moves with it (at the latest, December to February).

How to buy the team. The full three-person team is **£36,125 per month**. That rate is the same whichever deliverables TXN selects: you are buying the team for the period, not a fixed scope. Phase 3 is a committed three months, November to January, **£108,375 in total**, invoiced monthly in advance.

B · Phase 3 deliverables (November to January)

The remaining deliverables after the October launch, by component, worked across the three-month Phase 3 to January. This is a menu: TXN picks what is in scope and in what order within the window. Sizes (S to XL) are indicative effort only.

Agent Access Layer + A2A endpoint L

- Complete the tool catalogue from the finalised Core API docs, wired onto the live API once multi-tenancy and the move to TXN's Azure Europe are done.
- Product webhooks for change-impact triggers, which land late in DT's order, after Visa certification.
- Productionise the A2A endpoint for clients' own external agents, and finalise the audit store schema, retention, and PII handling with DT.

Internal Ops Agents, full suite XL

- Extend onboarding beyond v1: full CIQ assembly, scheme project and BIN handling, and project-plan automation.
- Settlement Reconciliation (promoted from Internal Ops on 18 June, dual-use internal and client-facing), after a data-capture session with DT.
- Scale the knowledge engine: proactive mining on a schedule, self-healing docs in production, and deeper CRM automation across the customer lifecycle.

Co-pilot M

- Impact preview before a change ("drop the limit from £200 to £100 and roughly 20% of recent transactions would decline"), powered by the data lake.
- The stuck-user replay mechanism, once Stackworkz exposes Console instrumentation after the 9 July handover.
- Deeper Console awareness and multi-page execution.
- Data-lake impact analysis for merged journeys.

10

B · Phase 3 deliverables, continued

The remaining components. Fraud & Risk Assist and Reconciliation are net-new, not started in the build phase; the rest deepen from their launch v1 to full version one.

Agent Inbox & advanced Alerts L

- The Agent Inbox: a central queue where proactive items and agent-proposed plans land for review and approval.
- Approvable plans, where an alert implies an action: the agent proposes it and executes on confirmation through the approval queue.
- Notification routing to the user's preferred channel: in-Console, Slack, Teams, or email.
- Data-lake anomaly detection at scale, plus cross-program benchmarking (needs multiple clients, contractual consent, and de-anonymisation safeguards).
- Finalise the central alert store and the observability integration with DT.

Developer Support s

- Sign-up and login, with the full four-level access gating deferred from the July portal.
- An autogenerated LLMS.txt-style downloadable doc for external agents.
- Voice support evaluation and a deeper lead-generation funnel.

Fraud & Risk Assist v1 M net-new

- Real-time enrichment and rules that advise, not decide.
- Depends on the data lake and live transaction volume (the data flywheel); requirements still being collected.

Reconciliation TBC net-new

- Auth and clearing settlement reconciliation, dual-use internal and client-facing.
- Scoped after a capture session with DT; the data format is not yet known.

A three-month phase, scope agreed together. Phase 3 runs November to January. The team works deliverables in parallel, so a substantial share of this menu is deliverable within the three months; we agree the selection and order together. The cost is £36,125 per month for the three months, £108,375 in total. If the start slips, the window shifts with it (at the latest, December to February).

Investment, continued

C · Payment schedule

Tranche	Trigger	Amount
Build M1	On commencement	£36,125
Build M2	Month 2	£36,125
Build M3	Month 3	£36,125
Build M4	Month 4 (launch)	£36,125
Build phase total		£144,500
Phase 3 M5	Month 5 (November)	£36,125
Phase 3 M6	Month 6 (December)	£36,125
Phase 3 M7	Month 7 (January)	£36,125
Phase 3 total · 3 months		£108,375

Payment schedule for the build phase to launch and the three-month Phase 3 that follows, all invoiced monthly in advance. If kickoff moves, the milestone months shift with it (Phase 3 at the latest December to February); the amounts do not change.

Investment, continued

D · Support & Maintenance fee

After the build and the three Phase 3 months, a fixed monthly support and maintenance fee covers the run, monitoring, evaluation upkeep, and the support and continuous-improvement effort that keeps the agentic layer current as the platform evolves. AI model consumption is **not charged by Novosapient**: it runs on TXN's own provider accounts.

SUPPORT & MAINTENANCE	MONTHLY
Managed support & maintenance (fixed)	£4,250
Indicative annual value	£51,000

Includes Agent Ops monitoring (model usage, cost tracking, agent performance, anomaly detection across the fleet), up to 20 hours of support and enhancement per month, and the service levels set out below. AI consumption is billed to TXN directly by the provider, not by Novosapient.

E · Total committed value

Sprint Zero (delivered)	£17,500
Build phase to launch (4 months)	£144,500
Phase 3 team, November to January (3 months)	£108,375
Total committed to version one	£270,375
Support & Maintenance, after Phase 3	£4,250 / mo

Total committed cost: £270,375, comprising Sprint Zero, the four build months to the October launch, and the three Phase 3 months to January at the unchanged £36,125 monthly rate. The Phase 3 scope is agreed from section B. Support and maintenance then runs at £4,250 per month.

Service levels and support

The £4,250 monthly fee is a managed service, not a break-fix contract. It runs the agentic layer, keeps it healthy as TXN's API and documentation change, and responds to issues against defined targets.

What the fee covers

Run & monitor

The Agent Ops platform watches model usage, cost, agent performance, and anomalies across the full fleet, with proactive alerting before issues reach users.

Keep current

Evaluation and regression upkeep as the Core API, data lake schema, and documentation change. The self-healing knowledge loop keeps answers accurate.

Support & improve

Up to 20 hours per month of support and continuous improvement: prompt tuning, guardrail adjustments, and small enhancements.

Incident response

Triage, workaround, and resolution against the severity targets below, with a single named point of contact.

Service level targets

SEVERITY	DEFINITION	RESPONSE	RESOLUTION TARGET
P1 Critical	Agentic layer unavailable across a surface, or an agent taking unsafe actions. Kill switch engaged.	15 minutes	Workaround within 4 hours
P2 High	A major component degraded for multiple clients with no workaround (e.g. Co-pilot or Inbox failing).	4 hours	1 business day
P3 Medium	A single feature impaired; a workaround exists.	1 business day	5 business days
P4 Low	Minor or cosmetic issue, or an enhancement request.	2 business days	Next release cycle

Service hours

09:00 to 17:00 UK, Monday to Friday, excluding public holidays. Response and resolution targets apply within service hours; outside them, P1 incidents are handled on a best-effort basis.

Change requests

Work beyond the monthly support allocation, or any new scope, is scoped as a change request and charged at our standard day rates (see section 9).

Commercial terms and acceptance

Validity	14 days from issue.
VAT	All amounts exclusive of VAT.
Build payment	Four monthly tranches of £36,125, first on commencement.
Phase 3 retainer	£36,125 per month for the committed three-month Phase 3, November to January (shifts with kickoff, at latest December to February), to deliver the agreed Phase 3 scope (section 10B). Invoiced monthly in advance.
Support & Maintenance	Fixed fee of £4,250 per month, commencing once Phase 3 ends, against the service levels in section 10. Invoiced monthly in advance. Change requests beyond the monthly allocation are charged at our standard day rates (section 10).
AI consumption & infrastructure	AI model usage runs on TXN's own provider accounts and is not charged by Novosapien. Cloud, vector store, and observability hosting are excluded, provisioned in TXN's own multi-region Azure (Europe) or by Novosapien with permission.
IP	TXN owns all intellectual property in the Deliverables: the components, agents, source code, models, prompts, configurations, and documentation Novosapien builds to meet the requirements agreed in Sprint Zero (the user stories, component and sub-component specifications, and acceptance criteria recorded in the TXN requirements set). Novosapien assigns that work product to TXN on creation, not contingent on final payment, and hands over the complete source, so TXN can operate, modify, and evolve the capability independently.
Late payment	Interest at 4% above the Bank of England base rate.
Liability	Capped at total fees paid. No indirect or consequential damages, except confidentiality or wilful misconduct.
Confidentiality	Two-year survival post-termination.
Termination	30 days' notice. Deliverables handed over in a deployable, documented state.
Warranties	Reasonable skill and care. The architecture is designed to support and not impede TXN's PCI-DSS, GDPR, and FCA obligations; TXN owns its certifications and compliance frameworks.
Governing law	England and Wales.

THE NEXT STEP

Ship the difference before the window closes.

Countersign to commence Phase 1. A working agentic pilot in six weeks, the launch-critical layer wired into production for October, then a three-month Phase 3, November to January, taking the remaining deliverables to version one on a scope agreed with you.